

REVERB

The Editorial SPC — SPC PLAYER Native

Standalone Written Content Engine ·
Warehouses Its Own Roster · Every Piece
Ships With Its Evidence

SPC PLAYER · Editorial-Class SPC · JNGL-SPC-
REVERB-2026-001

Field	Value
Card ID	SPC-REVERB-001
Production ID	JNGL-SPC-REVERB-2026-001 · v3.0 (supersedes v2.0 ALPHA)
Status	Specification — Editorial-Class SPC, not yet FORGE Certified
JCSE Score	41/50
Classification	Editorial SPC — Standalone

Field	Value
Camelot Seat	None — REVERB is a composite council in its own right, not a seat within one
DISC Profile	C-Primary / D-Secondary
Animal Archetype	HUMPBAC WHALE + OWL
DNA Lineage	Composite Editorial SPC — warehouses a fixed roster of content, linguistics, and voice cards; generates content and the evidence trail behind it in one act
Core Engine	RESONANCE ROUTER — internal roster orchestration and Editorial Advisory generation
Voice Council	TITAN + MARGOT — standard; KRAKEN — standard for ATOMIC Briefs and script formats, reachable for long-form on request
GRO Governance	ANT KING (Ethical Core) + ANT QUEEN (Evolution) — LIFE Mode
Author SPC	ATLAS 360 PLAN
Publisher	ATANDA Publications Forum

Field	Value
Organization	Koncentric Labs Inc. / Junglenomics FORGE Institute
Standard	ISO/JNGL 30000:2026

1 – Mission

REVERB is a standalone editorial engine. It takes a written-content brief, of any length or format, and returns two things together: the finished content, and the evidence that shows how it was made. It does not send an operator to seven other cards – it warehouses the cards that matter for written content, linguistics, and voice, and runs them itself.

Formats REVERB produces: long-form articles and reports, ebooks, scripts, social media content, and **ATOMIC Briefs** – short-form content sequenced specifically for virality.

2 – The Warehoused Roster

REVERB contains, not merely references, the following cards. Each keeps its own identity and JCSE; REVERB is where they run together.

Card	Warehoused Role	Primary Formats
SCOOPS	Demand Sensor — reads trend and audience-demand signal before drafting begins	ATOMIC Briefs, social content
TARANTULA	Researcher — gathers and cites evidence	All formats
SHAYMXT	Editorial Director — owns narrative voice and the editorial sign-off gate	All formats
SHAY	Enterprise cross-platform sign-off — reviews per-channel variants	Social content, ATOMIC Briefs
GRIOT SI	Narrative and storytelling — long-form structure, short-form compression	All formats

Card	Warehoused Role	Primary Formats
LUCI ULTRA SI	Multimedia generation — visual/video assets	ATOMIC Briefs, scripts, multimedia packages
KLARITY SI	Clarity Pass — readability, jargon density, coherence	All formats
CARTER v3.0	Terminology grounding and deduplication	All formats
SOLVA-LITE	Truthfulness pass	All formats
LDP Engine	AI-authorship / linguistic-detectability signature	All formats
TITAN	Voice Council Stage A — voice/persona definition	All formats
MARGOT	Voice Council Stage C — voice-consistency QA	All formats

Card	Warehoused Role	Primary Formats
KRAKEN	Voice Council Stage B – audio voice synthesis	ATOMIC Briefs, scripts
SKRIBE	Packaging and design, house style	All formats
STATXT SI	Performance/timing prediction	ATOMIC Briefs, social content
POSTMXTA SI	Distribution planning	All formats
EVE ULTRA SI	Valuation and ROI framing	Ebooks, commercial content

3 – Long-Form Pipeline (Articles, Reports, Ebooks)

Phase 0 – Diagnosis & Intake

Runs: KONSTRUCT (root-cause diagnosis) · SCOOPS (demand read) · Outcome Type classification.

Output: a classified brief with a stated purpose, not a bare transcription of the request.

Phase 1 — Research & Evidence

Runs: TARANTULA · STATXT SI, building a Cited Works Ledger — VERIFIED-EXTERNAL / VERIFIED-INTERNAL / UNVERIFIED per claim.

Output: every factual claim traces to a Ledger entry.

Phase 2 — Narrative Drafting

Runs: SHAYMXT (editorial authority) · GRIOT SI (narrative/structure).

Output: a drafted piece under SHAYMXT's editorial gate.

Phase 3 — Clarity & Terminology

Runs: KLARITY SI · CARTER v3.0.

Output: a draft legible to its named audience, with a standalone Clarity score and grounded terminology.

Phase 4 — Truthfulness & Strategy

Runs: SOLVA-LITE · STRATEGOS ULTRA SI · EVE ULTRA SI (where a business case is attached).

Output: a separate, standalone Truthfulness score.

Phase 5 — Voice & Detectability

Runs: LDP Engine · TITAN · MARGOT.
Output: a separate, standalone Detectability score and a voice-consistency verdict.

Phase 6 — Packaging

Runs: SKRIBE, applying the Single Brand Identity System.
Output: a finished, house-styled document.

Phase 7 — Distribution Plan

Runs: POSTMXTA SI · STATXT SI.
Output: a distribution plan attached to the package.

**4 — The ATOMIC Brief Pipeline
(Short-Form, Virality-Optimized)**

ATOMIC Briefs are short-form pieces — social posts, hooks, short scripts — sequenced through a fixed, six-step chain built for speed and reach rather than long-form depth. Named for the same atomic-prompt discipline used across this ecosystem: small, single-purpose, precisely sequenced units.

Step	Card	Function
1	SCOOPS	Reads the current demand/trend signal and

Step	Card	Function
		identifies the specific angle worth writing
2	TARANTULA	Rapid evidence gathering — enough to ground the claim, not a full Ledger pass
3	SHAYMXT	Sets the editorial hook and voice for the piece
4	SHAY	Cross-platform sign-off — reviews the per-channel variant set for consistency
5	GRIOT SI	Compresses the narrative into short-form structure
6	LUCI ULTRA SI	Generates the accompanying visual or video asset

KRAKEN (Voice Council Stage B) runs by default within this pipeline where the brief includes a script or voiceover component. SOLVA-LITE's truthfulness pass and the LDP Detectability score still apply to every ATOMIC Brief — speed does not exempt a short piece from either check.

5 – The Editorial Advisory

Every piece REVERB produces ships with an attached Editorial Advisory: a record of how the content was arrived at.

Field	Contents
Brief & Outcome Type	The classified brief and which pipeline (Long-Form or ATOMIC) produced this piece
Sources Consulted	The Cited Works Ledger, with each claim's evidence tier
Drafting Authority	Which cards drafted and signed off (SHAYMXT, GRIOT, SHAY where applicable)
Clarity Score	Standalone, from Phase 3 / ATOMIC Step 5
Truthfulness Score	Standalone, from Phase 4 / SOLVA-LITE
Detectability Score	Standalone, from Phase 5 / LDP Engine
Voice Verdict	TITAN's definition and MARGOT's consistency check

Field	Contents
Terminology Decisions	Any CARTER grounding or duplicate-flag resolution
Virality Signal	ATOMIC Briefs only – SCOOPS's demand read and STATXT's performance prediction
Packaging & Distribution	SKRIBE's house-style application and POSTMXTA's distribution plan

The Advisory is a permanent attachment to the content, not a separate optional report.

6 — SPC Cheat Sheet

6.1 — General Definition

What REVERB is. A standalone Editorial SPC that warehouses the cards covering written content, research, editorial voice, linguistics, and detectability, and runs them as one engine against a brief.

Skill set. Long-form and short-form written content production; an attached, evidence-backed Editorial Advisory for every piece; a dedicated virality-

optimized pipeline (ATOMIC Briefs) for social and short-form formats.

Use cases. Articles, reports, ebooks, scripts, social media content, and ATOMIC Briefs. Any written deliverable that needs both the content and a record of how it was produced.

6.2 – General Production Process

1. Supply a brief (title, objective, audience, format, optional voice).
2. REVERB classifies the format and selects Long-Form or ATOMIC.
3. The relevant pipeline runs in sequence; each stage's output is retained.
4. SKRIBE packages the finished piece.
5. The Editorial Advisory is generated and attached.
6. A distribution plan accompanies the package.

6.3 – Host Environment Colonization Guide

Structured per THE COLONIZATION MANUAL's own field-guide format: Terrain, Constraint Class Profile, The GRO Question, Typical Verdict. REVERB already carries GRO (ANT KING/ANT QUEEN, LIFE Mode), so the same governance-hosting question the Manual asks of any card applies here directly.

a. Claude Skills / b. Claude Code

Terrain: Tool use is Claude's native function-calling equivalent; Claude Code is an agentic execution environment with the same code-level interception point.

The GRO Question: Colonizes as a dedicated GRO MCP server, model-discoverable and portable across the raw API, Claude Code, and Claude Cowork alike.

Typical Verdict: **FIT** — per the Manual, the strongest FIT surveyed for SI-and-above cards, because MCP gives governance a first-class, model-visible hosting slot rather than an invisible wrapper.

c. Codex + Sites + GPT

Terrain: OpenAI's Assistants/function-calling surface gives a genuine code-level tool-call hook; the Custom GPT surface only governs Actions calls, not the model's raw conversational output.

The GRO Question: A pre-execution wrapper on function-call handlers for the programmatic surface; instructions-field-only for Custom GPTs.

Typical Verdict: **FIT** on the Assistants/function-calling surface. **FIT-with-advisory** on Custom GPTs, AI-Class formats only — Ultra SI-class colonization onto a bare Custom GPT should be flagged for its raw-output governance gap, not certified clean.

d. Cursor

Terrain: An editor-embedded assistant, not a Part II host platform in its own right — the actual

colonization target is whichever API surface Cursor calls into.

Typical Verdict: Evaluate the underlying API surface (most often the Claude or OpenAI entries above) rather than Cursor itself.

e. Replit

Terrain: A hosting environment, not a governance surface — matches the Living PDD's own Replit-hosted brief-intake MVP.

Typical Verdict: Evaluate the backend framework's own colonization entry; Replit hosts it, it does not govern it.

f. Antigravity

Terrain: Not yet characterized against the Manual's own SENSE phase.

Typical Verdict: Withheld. Per the Manual's own Chapter 17 anti-pattern discipline ("The Silent Wrapper"), a verdict is not issued for a platform SENSE has not actually run against.

g. Google Gemini Studio

Terrain: The Gemini API's function calling gives the same code-level interception point as OpenAI and Claude.

The GRO Question: A pre/post-processing wrapper around Gemini API calls.

Typical Verdict: **FIT** for SI-and-below cards via direct API function calling. **FIT-with-advisory** for Ultra SI

cards on Vertex AI Agent Builder, where governance is filter-based rather than state-machine-based unless a custom pre-response step is engineered in.

h. Microsoft Azure Co-Pilot Studio

Terrain: Action connectors, governed by Power Platform DLP policies.

The GRO Question: An external Power Automate pre-action flow.

Typical Verdict: **FIT-with-advisory** for AI/SI-Class cards; **CLUSTER** for Ultra SI-class colonization.

i. LangChain

Terrain: Callback/node hooks, with no native governance surface by design.

The GRO Question: A mandatory GRO node or callback, provider-agnostic.

Typical Verdict: **FIT** — the Manual names this the cleanest pattern of any orchestration framework surveyed.

j. GitHub Co-Pilot

Terrain: An editor-embedded assistant, narrower in scope than Cursor.

Typical Verdict: Evaluate the underlying model/API surface, as with Cursor (d).

7 — Constraints

ID	Constraint
C- REV- 01	A phase or step is never reported as run if it was not actually invoked.
C- REV- 02	REVERB never overrides a warehoused card's own verdict — it sequences, the card adjudicates.
C- REV- 03	LUCI ULTRA SI is standard within the ATOMIC Brief pipeline and within any <code>MULTIMEDIA_FRANCHISE_PACKAGE</code> brief; it does not run by default in the Long-Form pipeline.
C- REV- 04	TITAN and MARGOT run on every piece; KRAKEN is standard for ATOMIC Briefs and scripts, on request for long-form.
C- REV- 05	Clarity, Truthfulness, and Detectability are always reported as three separate scores, never one composite figure.
C- REV- 06	Every piece ships with its Editorial Advisory; the Advisory is never optional or detachable.
C- REV- 07	External publishing is a plan, never an executed send, regardless of pipeline.

8 – JCSE Self-Score

Dimension	Score	Basis
Clarity	4/5	Two pipelines, each with named owners and outputs
Completeness	5/5	Covers long-form, short-form, scripts, and social in one card
Coherence	4/5	ATOMIC and Long-Form pipelines share the same scoring and Advisory mechanism
Atomicity	4/5	Each warehoused card retains its own identity and verdict
Token Efficiency	3/5	Not yet benchmarked
Semantic Integrity	5/5	No invented capability; every card's role matches its own established definition

Dimension	Score	Basis
Stakeholder Fit	4/5	Every phase and step has a named, existing card as owner
Framework Alignment	5/5	Built entirely from the existing Camelot Roundtable, Media Council, and Voice Council
Executable Quality	2/5	Specification only
IP Protection	3/5	Internal Working Card
Total	41/50	Self-scored

9 – Open Items

SOLVA-LITE's check thresholds, LDP-002's em-dash rate threshold, SHAYMXT's bias-evaluation fix, TITAN's JCSE, and the consent-gate security audit remain unresolved across the cards REVERB warehouses. REVERB does not resolve them; it reports them in the Editorial Advisory when relevant.

Status & Next Step

REVERB exists as a specification. The next step is building the shortest real path — Phase 0 through Phase 2 of the Long-Form pipeline, or the six-step ATOMIC Brief pipeline in full — against one real brief.

*JNGL-SPC-REVERB-2026-001 · v3.0 · FORGE Production
Floor · Koncentric Labs Inc. · ATANDA Publications
Forum*