

ATANDA Command Centre — ATLAS 360 Living PDD

Product, Project & Program synthesis of the current codebase

Production ID: `JNGL-ACC-PDD-CUR-2026-002`

Version: `2.0.0-draft`

Snapshot: `2026-09-03T13:11:24-05:00` (America/Chicago)

Source method: ATLAS 360 PLAN v3.0.0, Living PDD Engine

Supersedes: `JNGL-ACC-PDD-CUR-2026-001` (30 May 2026)

Part 0 — Title, Contents, Provenance & Honesty Gate

0.1 Document identity

Field	Current value
Human principal	Oluseye Shay Amusa
Document title	ATANDA Command Centre — ATLAS 360 Living PDD
Production ID	`JNGL-ACC-PDD-CUR-2026-002`
Date/time stamp	`2026-09-03T13:11:24-05:00`
Mode	ATLAS 360 PLAN — Living PDD Engine
Path	Path A (traced), followed by re-baselining because the earlier PDD is materially stale
GRO mode	LIFE MODE
Honesty tier	Tier 2 — implemented and internally test-verified; current public-production health was not re-certified in this pass
Certification	Draft architecture-of-record; no new JCSE, HIVE, MATHMON, SOLVA or FORGE certification claimed
Review cadence	Refresh after any merged capability, access-model, integration, schema, engine or deployment change

0.2 Contents

- 0. Title, Contents, Provenance & Honesty Gate
- 1. Statement of Requirements

- Business Case, Financial Position, Demographics & Validation
3. SPC Roster, Skills & Contributory Roles
 4. Platform Specifics
 5. Atomic Task Worksheet
 6. DJ Implementation & Operations Plan
 7. Summary of Sources
 8. Definition of Terms
 9. IP Pre-Registration Screening

0.3 What this document claims

- It reverse-translates the repository and running development workflows as they existed at the snapshot time.
- It distinguishes active behavior from dormant code and design-intent.
- Its implementation counts are code-derived: 22 distinct frontend paths, 139 Express route declarations, 33 Drizzle table declarations, 25 mounted route modules, and 15 POST-based HARNESS execution paths plus the escalation stream.
- It records the latest automated evidence available in the workspace: full TypeScript typecheck passed; API tests passed 270/270; web tests passed 94/94; all 11 committed database migrations were applied in development.

0.4 What this document does not claim

- It does not claim that implemented subscription, marketplace, IP-search or external-service seams are commercially live.
- It does not claim revenue, ROI, valuation, market size, adoption or user-demographic facts that cannot be derived from code.
- It does not claim that a passing development test suite proves the latest published deployment is healthy.
- It does not claim legal, regulatory, security or IP clearance.
- It does not reclassify HARNESS engines as SPCs. Engines are the instruction layer; the artifacts they produce may be SPCs or PDDs.

0.5 OSIRIS drift finding

The superseded PDD described an active public subscription portal with 101 endpoints, 20 tables and 13 engines. The current product is an internal staff operations platform, while Clerk/Stripe subscription behavior is dormant behind [SUBSCRIPTIONS_ENABLED](#). The current repository has 139 route declarations and 33 table declarations, plus new F0/F0.5/MAP, acquisition magnets, cost-cap operations and cron-recovery capabilities.

Drift band: RED. Route growth is approximately 38%, table growth is 65%, and the active access/commercial posture changed. Under ATLAS 360 PLAN §11, the prior baseline is therefore superseded rather than presented as current or FORGE-verified. This document establishes the new baseline.

Part 1 — Statement of Requirements

1.1 Systems diagnosis

The present system solves an operational translation problem: staff need one controlled workspace for taking an idea or existing specification through a repeatable PromptWare production pipeline, preserving evidence, costs, audit history and exportable outputs.

The architecture must support two postures without rebuilding:

1. **Active posture — internal staff operations.** A shared access code gates the app; typed names/initials provide attribution; authenticated staff receive ADMIN/INSTITUTION-equivalent access.
2. **Dormant posture — subscription SaaS.** Clerk identity, Stripe subscriptions, organizations, seats, per-project credits and F1000 remain implemented and tested, but are disabled by default with `SUBSCRIPTIONS_ENABLED=false`.

1.2 Current functional requirements

ID	Priority	Requirement	Status
ACC-001	MUST	Run the product as a pnpm/TypeScript monorepo with separately routable web and API artifacts	PRESENT
ACC-002	MUST	Authenticate active staff with the shared access code and a signed session cookie	PRESENT
ACC-003	MUST	Attribute staff actions to a provisioned local user without exposing the shared code	PRESENT
ACC-004	MUST	Preserve dormant Clerk/Stripe SaaS code behind one reversible feature flag	PRESENT
ACC-010	MUST	Move a session through the F1!F7 PromptWare sequence and persist each artifact	PRESENT
ACC-011	MUST	Require a SPARTAN-certified MVP-PDD/PWDD before F8 Code DJ generates a code bundle	PRESENT
ACC-012	MUST	Provide F0 advisory/retainer management and monitoring outside the core F1!F8 path	PRESENT
ACC-013	MUST	Support F0.5 diagnostic intake and MAP/Mathmon structured analysis paths	PRESENT

ACC-014	SHOULD	Provide ATLAS-J crystallisation, DE-SPC synthesis, VIBE DJ, HOST DJ and PFP drift analysis as side-steps	PRESENT
ACC-020	MUST	Treat OpenAPI as the primary API contract and generate client hooks/ Zod schemas where supported	PRESENT WITH DISCLOSED EXCEPTIONS
ACC-021	MUST	Authorize user-owned data at the application layer on every query	PRESENT
ACC-022	MUST	Gate every LLM execution path with the active company-wide monthly cost budget	PRESENT
ACC-023	MUST	Record provider, token and USD telemetry for LLM runs, including pre-session work	PRESENT
ACC-024	MUST	Stream long-running engine finalization that can exceed proxy request limits	PRESENT
ACC-030	SHOULD	Accept manual sessions, ingested IPDDs and Advanced Cartridge project context	PRESENT
ACC-031	SHOULD	Export certified work as files, IDE bundles and GitHub-ready repositories	PRESENT
ACC-032	SHOULD	Provide public certificate verification and acquisition magnet tools	PRESENT
ACC-033	SHOULD	Compute badges and Ascension progress from persisted evidence rather than self-report	PRESENT
ACC-040	MUST	Apply committed database migrations without destructive schema pushes	PRESENT
ACC-041	MUST	Keep Stripe webhook processing raw-body-first and idempotent when subscriptions are enabled	PRESENT, DORMANT
ACC-042	MUST	Send operational email safely, escape dynamic HTML and never stamp failed sends as delivered	PRESENT

ACC-043	MUST	Monitor scheduled jobs for stale, flaky and recovered states	PRESENT
ACC-044	SHOULD	Expose staff cost, activity and operations-health dashboards	PRESENT
ACC-045	SHOULD	Mirror local changes to GitHub without force-pushing or requiring workflow scope	PRESENT, CURRENTLY NEEDS BASE RECONCILIATION

1.3 Non-functional requirements

- **Security:** signed sessions, no secret values in outputs, SSRF controls on URL evidence, app-layer ownership checks, sanitized email HTML and server-recomputed safety gates.
- **Reliability:** Stripe replay protection, exactly-once alert stamps where chosen, migration-status startup checks, cron heartbeat history and recovery notification.
- **Cost control:** every LLM callsite must be budget-gated and recorded; no silent unmetered pre-session calls.
- **Portability:** certified artifacts can be exported as Markdown/PDF/ZIP/IDE/GitHub bundles.
- **Maintainability:** OpenAPI-first where practical, generated clients, one schema module per domain, committed migrations and workspace-wide typecheck/tests.
- **Observability:** structured API logging, engine run ledger, cost views, activity views, cron status and optional Sentry.

1.4 Known requirement tensions

- Contract-first is the default, but org/activity/cartridge and other specialist routes use hand-written validation because multipart/OpenAPI codegen has known limitations.
- Application-layer authorization has no PostgreSQL RLS safety net; every query must retain explicit ownership filters.
- Staff mode intentionally bypasses subscription tier and badge gates, while keeping those systems in the codebase.
- Cartridge injects authoritative context into every call, but the operator still manually drives F1!F7.
- Marketplace and external Ark/Sphinx seams fail explicitly when not configured; they are not mocked as live.

Part 2 — Business Case, Financial Position, Demographics & Validation

***ATLAS 360 Path A honesty rule:** business rationale not provable from code is marked `[INFERRED]`. No SOLVA GAUNTLET was run for this draft.*

2.1 Business case

`[INFERRED]` ATANDA Command Centre productizes a repeatable method for turning ideas and source documents into structured PromptWare deliverables, reducing dependence on untracked one-off prompting. The current staff posture supports internal production and quality control before any broader subscription relaunch.

2.2 Value already embodied in the code

Capability	Operational value	Evidence class
Persistent F1!F8 session pipeline	Reproducible work with a traceable artifact chain	VERIFIED-INTERNAL
PFP and server-recomputed verdicts	Prevents a model from self-approving contradictory drift results	VERIFIED-INTERNAL
Company-wide LLM cap and run ledger	Makes spend visible and mechanically bounded	VERIFIED-INTERNAL
F8 ZIP/IDE/GitHub delivery	Converts a certified specification into a portable build handoff	VERIFIED-INTERNAL
IPDD ingestion and Cartridge	Brings existing project material into the same controlled pipeline	VERIFIED-INTERNAL
F0 retainers and monitoring	Supports ongoing advisory operations around the core build flow	VERIFIED-INTERNAL
Scheduled-job health and recovery email	Makes silent cron failures and recoveries visible to operators	VERIFIED-INTERNAL

2.3 Financial position

- Real LLM cost is recorded per run and summarized by user, engine and month.
- The active staff guardrail is [STAFF_MONTHLY_COST_CAP_USD](#).
- Stripe subscription, seat and one-time purchase code is implemented but dormant.
- No revenue, margin, payback-period, valuation or savings claim is made in this document.
- Any future financial forecast requires EVE modelling and a SOLVA GAUNTLET verdict before publication.

2.4 Users and stakeholders

Group	Present interaction
Staff operator	Creates sessions, runs engines, ingests context, exports outputs and reviews personal cost/activity
Staff administrator	Reviews operational health, badges, cost alerts and scheduled-job history
Advisory operator	Manages F0 engagements, retainers, reports, tasks and monitoring alerts
Public verifier	Checks a certificate without entering the authenticated portal
Acquisition visitor	Uses the agent-test and savings-calculator magnets
Future subscriber/org member	Code exists but the experience is dormant until subscriptions are enabled

No demographic distribution or adoption count is asserted.

2.5 Validation status

Check	Snapshot result
Workspace TypeScript typecheck	PASS
API automated tests	270/270 PASS
Command Centre automated tests	94/94 PASS
Development API startup	PASS; server listening on port 8080
Development web startup	PASS; Vite listening on port 21731
Database migration status	11/11 committed migrations applied in development
Published production health	NOT RE-CERTIFIED IN THIS PASS
GitHub mirror	DEGRADED; last sync base is not an ancestor of current HEAD and needs explicit base reconciliation

Part 3 — SPC Roster, Skills & Contributory Roles

These are methodology overlays for this document. They are not runtime agents and do not reclassify HARNESS engines.

Contributor	Role in this Living PDD	Status applied here
ATLAS 360 PLAN	Host structure: product, project and program synthesis	Reference specification supplied by the principal
KONSTRUCT	Requirements diagnosis before implementation description	Applied analytically
KLARITY	Plain-language and audience separation	Applied analytically
OSIRIS	Baseline comparison and GREEN/AMBER/RED drift ruling	Applied analytically
CODE DJ	Reverse-engineering code into platform and atomic-task descriptions	Applied analytically
BUGMXT	Code/PDD fidelity and defect-risk framing	Applied analytically
SKRIBE	Living document consistency and readable tables	Applied analytically
CARTER	Terminology normalization	Applied analytically; no persistent terminology-store write claimed
SOLVA	Business-case adjudication	NOT RUN; Part 2 remains inferred
EVE	Financial modelling	NOT RUN

KONSA	External evidence retrieval	NOT RUN; Part 7 uses internal primary sources
ALBERT	IP candidate extraction and open-register screening	Candidate extraction only; no register search run

Runtime PromptWare roster

The implemented instruction layer consists of:

- Core path: F1, F2, F3, F4, F5, F6, F7 and F8 Code DJ.
- Advisory/diagnostic paths: F0, F0.5 and MAP.
- Side-steps: F6 VIBE DJ, F8 HOST DJ, DE-SPC evolve, ATLAS-J crystallisation and PFP.
- Public acquisition engines: agent test kit and savings calculator.

The source directory also contains shared prompt, routing and utility modules; file count is not used as an engine count.

Part 4 — Platform Specifics

4.1 Product topology

Layer	Present implementation
Web artifact	React + Vite + Wouter; 22 distinct route paths; public and authenticated route sets
API artifact	Express 5; 25 mounted route modules; 139 route declarations
Data	PostgreSQL + Drizzle; 33 `pgTable` declarations; 11 committed migrations
API contract	`lib/api-spec/openapi.yaml` with Orval-generated Zod and React Query packages
AI providers	Provider-switching integration packages for Anthropic, OpenAI and Gemini; prompts use recorded provider/ model metadata
Documents	Markdown/PDF/export helpers plus IDE and repository bundles
Replit services	Static web artifact, runnable API artifact and isolated mockup-sandbox design artifact

4.2 Active navigation and user journeys

Public paths:

- [/verify](#)
- [/test-your-agent](#)
- [/calculate-your-savings](#)

Authenticated paths:

- [/command](#), [/guide](#), [/sessions](#), [/session/new](#), [/session/:id](#)
- [/ingest](#), [/cartridge](#), [/f0](#), [/prompts](#)
- [/quests](#), [/ascension](#), [/account](#)
- [/admin/badges](#), [/admin/ops](#)
- [/me/activity](#), [/me/costs](#)
- [/exemplars](#), [/exemplars/:id](#)

4.3 Core session lifecycle

1. Create or seed a session manually, from an IPDD, from an exemplar, or from Cartridge context.
2. F1 diagnoses the raw prompt against Context Craft dimensions.
3. F2 produces an Atomic Prompt.
4. F3 produces a CELL Micro Agent birth package.
5. F4 produces a Micro PDD.
6. F5 produces a full SPC.
7. F6 produces an ATLAS PDD and may invoke VIBE DJ.
8. F7 compresses and certifies an MVP-PDD/PWDD and exposes a verification path.
9. F8 Code DJ produces a bounded codebase bundle; GitHub/IDE export uses that same generated file set.
10. HOST DJ may provide advisory hosting ranking and deployment journey; it does not deploy.

4.4 Supporting domains

- **F0 operations:** engagements, reports, report codes, retainers, retainer tasks and monitoring runs.
- **Prompt quality:** F0.5 intake, MAP/Mathmon, DE-SPC, ATLAS-J and PFP.
- **Context intake:** ingestion credits/documents and Cartridge packages/documents/SPCs/links.
- **Progression:** badges, revocation history, Context Craft pillars, JST assessments, reader onboarding and Ascension.
- **Commercial layer:** subscribers, pricing content, F1000 invitations, organizations/members/invites and Stripe event ledger; dormant in staff mode.
- **Operations:** cost notifications, cron tick status/events, stale/flaky/recovery episodes, email preferences and integration credentials.
- **Growth:** exemplars, public verification and magnet sessions.

4.5 Data and authorization

- Drizzle schema modules are the database source; migrations are generated and committed.
- The app uses explicit user/org filters rather than RLS.
- Stripe webhook event IDs provide replay idempotency when billing is active.
- Staff sessions provision local user/subscriber records for attribution.
- Integration credentials are encrypted at rest where applicable; raw values are not returned to clients.

4.6 LLM execution and cost controls

- [harnessRoute](#) centralizes authentication, optional tier/rate gates, extra middleware, cost-budget enforcement and the engine handler.
- In active staff mode, the company-wide monthly cap supersedes personal tier caps.
- Every LLM run must create a run-ledger row; pre-session work uses nullable session IDs and

explicit sentinel engine IDs where required.

- Large finalization flows use streaming to avoid the proxy's single-request timeout.
- Tests use deterministic provider fixtures rather than live network calls by default.

4.7 Notifications and scheduled operations

Scheduled targets:

- Weekly activity digest
- Daily rate-limit reset
- Weekly F0 monitoring
- Cost-cap alert sweep every 15 minutes

Cron handlers record heartbeats. The monitor detects stale schedules, degraded/flaky frequency, and recovery. Exactly-once stamps are used to avoid alert floods. Email templates escape dynamic HTML, URL bases are restricted to HTTP(S), and callers check the email result before recording delivery where retryability matters.

4.8 Present limitations and design debt

- Some API routes intentionally remain outside OpenAPI codegen.
- No RLS fallback exists beneath application authorization.
- Subscription and organization billing journeys are not active in staff mode.
- Sphinx and Ark integration behavior depends on external services and configuration.
- GitHub mirror workflow currently needs a one-time explicit base reconciliation after history divergence.
- Public-production health is not declared current by this document.
- ALBERT register access and a live Part 9 search are not implemented as part of this repository.

Part 5 — Atomic Task Worksheet

The chain is **Feature !' Atomic Task !' Atomic Prompt**. Existing IDs are preserved; new rows extend rather than renumber the prior living PDD.

Feature	Task ID	Atomic prompt	Input	Verifiable output
Foundation	ACC-001	Start the API with migration status checked	env + schema	listening API or explicit startup failure
Foundation	ACC-002	Authenticate a staff operator	access code + display name	signed staff session
Foundation	ACC-003	Resolve the staff session to a local user	signed cookie	`req.localUser`
HARNESS	ACC-010	Diagnose one raw prompt	session + prompt	persisted F1 artifact
HARNESS	ACC-011	Build one Atomic Prompt	F1 artifact	persisted F2 artifact

HARNESS	ACC-012	Build one CELL package	F2 artifact	persisted F3 artifact
HARNESS	ACC-013	Build one Micro PDD	F3 artifact	persisted F4 artifact
HARNESS	ACC-014	Build one full SPC	F4 artifact	persisted F5 artifact
HARNESS	ACC-015	Draft one ATLAS PDD	F5 artifact	persisted F6 artifact
HARNESS	ACC-016	Certify one MVP-PDD/PWDD	F6 artifact	certificate + verification identity
Delivery	ACC-017	Generate one bounded codebase bundle	certified artifact	ZIP/manifest/file set
Fidelity	ACC-018	Compare one bundle against its certified PDD	PDD + bundle	server-normalized PFP report
Advisory	ACC-050	Create one F0 engagement	operator input	persisted engagement
Advisory	ACC-051	Sweep one eligible retainer	retainer + latest state	monitoring run + optional alert
Diagnostics	ACC-052	Run one F0.5 diagnostic intake	structured concern	diagnostic artifact
Diagnostics	ACC-053	Stream one MAP analysis	intake/context	persisted map
Intake	ACC-030	Claim and ingest one IPDD	document	source record + PWDD session
Intake	ACC-031	Claim and assemble one Cartridge	scoped project assets	package + linked session
Governance	ACC-040	Apply one generated schema migration	committed migration	matching DB journal
Governance	ACC-041	Record one LLM run cost	provider response	run-ledger row
Governance	ACC-042	Reject one over-budget LLM request	monthly spend + cap	HTTP 402 without new LLM call
Notifications	ACC-043	Dispatch one digest only after provider success	digest payload	sent count + delivery stamp
Notifications	ACC-044	Close one stale cron episode after recovery	new heartbeat	recovery stamp + all-clear email
Integration	ACC-045	Push one F8 file set to GitHub	bundle + connection	repository commit or explicit error

Growth	ACC-060	Record one acquisition-magnet session	public tool input	privacy-bounded magnet result
Progression	ACC-061	Award one evidence-backed badge	persisted qualifying signal	idempotent badge state

5.1 Worksheet acceptance rules

- One task performs one coherent operation.
- Every task traces to a feature and a current code path.
- LLM tasks include both budget enforcement and run recording.
- Any client-provided aggregate, score, verdict or ranking with gate impact is recomputed server-side.
- A task is not marked complete merely because an email promise resolved; [EmailResult.ok](#) must be true.

Part 6 — DJ Implementation & Operations Plan

6.1 As-built tool selection

Tool/platform	Current role
Replit	Project workspace, artifacts, workflows, secrets, development preview and publishing
pnpm workspaces	Dependency graph and package execution
TypeScript 5.9 / Node.js 24	Primary implementation/runtime
React + Vite	Staff web interface
Express 5	API and webhook service
PostgreSQL + Drizzle	Persistent data and migrations
OpenAPI + Orval	Contract and generated clients/validators
Vitest	API and web automated tests
GitHub integration	Repository mirror and F8 delivery target
Stripe	Dormant subscription, seat and credit billing
Replit AI integrations	Provider access for model execution

6.2 Change protocol

1. Diagnose the requirement before naming a new package or platform.
2. Update OpenAPI first unless the route is an explicitly documented exception.
3. Add/change Drizzle schema and generate a committed migration; never synchronize shared

DBs with destructive [push](#).

4. Use canonical [harnessRoute](#) ordering for new LLM execution surfaces.
5. Ensure every LLM caller is budget-gated and recorded.
6. Add focused tests and run workspace typecheck/test.
7. Restart affected workflows once after the coherent change.
8. Update this Living PDD when the capability, count, posture or invariant changes.

6.3 Operational baseline

- Web development workflow: running.
- API development workflow: running with all migrations applied.
- Mockup sandbox: running as an isolated design artifact.
- Typecheck and test workflows: passing at the snapshot.
- GitHub Sync: running but failing fast on a divergent resume base; do not force-push.

6.4 Next review triggers

Refresh this document immediately when any of the following occurs:

- [SUBSCRIPTIONS_ENABLED](#) becomes active.
- A new engine or public/authenticated route is added.
- The API contract exception list changes.
- A schema table or migration is added.
- A new scheduled target, alert stamp or integration is introduced.
- Published production health is formally reverified.
- ALBERT performs the first live register query for this product.

Part 7 — Summary of Sources

7.1 Internal primary sources

Source	Use	Classification
`replit.md`	Active posture, stack, operating constraints, scheduled jobs and terminology	VERIFIED-INTERNAL
`README.md`	Product summary and repository entry points	VERIFIED-INTERNAL
`docs/architecture/decisions.md`	Binding architecture decisions and known contract exceptions	VERIFIED-INTERNAL
`docs/architecture/features.md`	Detailed shipped feature behavior	VERIFIED-INTERNAL
`docs/architecture/gotchas.md`	Non-obvious implementation constraints	VERIFIED-INTERNAL
`docs/architecture/llm-cost-guardrail.md`	Cost-budget invariant	VERIFIED-INTERNAL

`lib/api-spec/openapi.yaml`	Contract-declared API surface	VERIFIED-INTERNAL
`lib/db/src/schema/*.ts`	Current data model	VERIFIED-INTERNAL
`artifacts/api-server/src/routes/*.ts`	Runtime API route declarations	VERIFIED-INTERNAL
`artifacts/api-server/src/engines/*.ts`	PromptWare execution layer	VERIFIED-INTERNAL
`artifacts/command-centre/src/App.tsx`	Active browser route map	VERIFIED-INTERNAL
Workspace workflow logs at snapshot	Typecheck, tests, startup and migration evidence	VERIFIED-INTERNAL

7.2 Method source

[ATLAS_360_PLAN_v3_0_0_ALBERT_Integration_1788458927280.pdf](#), supplied 3 September 2026, provides the 10-Part standard, Path A/Path B Living PDD rules, OSIRIS drift bands, source classification and ALBERT Part 9 boundary.

7.3 External evidence status

No KONSA evidence-retrieval pass was run. Business and market statements are therefore limited to code-supported operational value or marked [\[INFERRED\]](#).

Part 8 — Definition of Terms

Term	Type	Grounded definition in this project
ATANDA Command Centre	Product	Internal staff operations platform around the FORGE.BONSAI HARNESS
FORGE.BONSAI HARNESS	Governance/architecture	Ordered PromptWare instruction layer that drives session transformations; not itself an SPC
Engine	Runtime component	A bounded instruction/execution path such as F1, F7, F8 or PFP
SPC	Artifact/method	Super Prompt Card; may describe a reusable governed prompt method
PDD	Artifact family	PromptWare Design Document unless explicitly qualified otherwise
IPDD	Input artifact	Ingestion Product Design Document supplied to the ingestion flow
PWDD	Output artifact	PromptWare Design Document produced from an ingested IPDD

Micro PDD	Qualified output	F4 intermediate document
ATLAS PDD	Qualified output	F6 structured design document
MVP-PDD	Qualified output	F7 SPARTAN-compressed, certifiable output
Living PDD	Governance artifact	Current code-derived architecture baseline that is refreshed as implementation changes
Cartridge	Context package	Scoped project context containing documents, SPCs and non-secret links, injected into session prompts
PFP	Quality control	PDD Fidelity Protocol that compares a certified document to a code bundle
F0	Advisory layer	Engagement, report, retainer, task and monitoring domain outside the sequential build path
F0.5	Diagnostic layer	Structured pre-build diagnostic intake
MAP	Analysis path	Streaming structured analysis associated with Mathmon intake/ map records
Staff mode	Access posture	Shared-code authentication with ADMIN/INSTITUTION-equivalent access
Subscription mode	Dormant posture	Clerk/Stripe tiered SaaS behavior enabled by `SUBSCRIPTIONS_ENABLED`
Exactly-once stamp	Reliability construct	Unique persisted marker used to prevent duplicate operational notifications
Drift band	Governance metric	OSIRIS classification: GREEN under 15%, AMBER 15–30%, RED over 30% or governance breach
DESIGN-INTENT	Evidence class	Specified but not proven as a live integrated capability

CARTER persistent-library synchronization was not run; these definitions are grounded in current repository language only.

Part 9 — IP Pre-Registration Screening

Preliminary candidate inventory only. No USPTO, WIPO, EUIPO, UKIPO, US Copyright Office or other register was queried in this pass. Nothing below states that a mark is clear, safe, available, registrable or patentable. This is not legal advice or a substitute for a licensed practitioner's clearance search.

9.1 Candidate extraction

Candidate	Apparent class	Repository use	Search status
ATANDA Command Centre	Product/service name	Product title and interface	UNSEARCHED
FORGE.BONSAI HARNESS	Product/method name	Core PromptWare instruction layer	UNSEARCHED
ATLAS 360 PLAN	Method/document name	Living-PDD methodology supplied for this pass	UNSEARCHED
Advanced Cartridge	Feature/service name	Scoped project-context intake	UNSEARCHED
PDD Fidelity Protocol / PFP	Method name	Code-to-PDD drift analysis	UNSEARCHED
ATLAS-J / JSON Crystallisation	Feature/method name	Typed ATLAS PDD transformation	UNSEARCHED
HOST DJ / VIBE DJ / CODE DJ	Productized role names	Advisory and implementation paths	UNSEARCHED
Ascension Protocol	Program name	Evidence-based onboarding/progression	UNSEARCHED
SPHINX Marketplace	External product/integration name	Outbound listing seam	EXTERNAL OWNERSHIP/RIGHTS NOT ASSESSED
F0, F0.5 and MAP mechanisms	Technical workflow candidates	Advisory and diagnostic architecture	PATENT LANDSCAPE NOT RUN
Exactly-once cron stale/flaky/recovery model	Technical mechanism candidate	Operational alert lifecycle	PATENT LANDSCAPE NOT RUN
Certified PDD-to-code drift gate	Technical mechanism candidate	PFP/F8 delivery control	PATENT LANDSCAPE NOT RUN

9.2 Required next ALBERT pass

If formal screening is requested:

1. Confirm target jurisdictions and product/service classes.
2. Run dated [SEARCH](#) requests for proposed names.
3. Run [PATENT_VALIDATION](#) only as a landscape search for mechanism candidates.
4. Name every register, query, class and query date in the findings table.
5. Keep any suggested replacement names marked [UNSEARCHED](#) until separately queried.
6. Route active disputes, received legal correspondence, litigation or deadlines within five business days directly to qualified counsel.
7. Do not generate filing documents or legal correspondence by default.

Living Update Record

Version	Date	Change
1.0.0	2026-05-30	Original ATLAS 4-Part reverse-translation baseline
2.0.0-draft	2026-09-03	Re-baselined under ATLAS 360 PLAN v3.0.0; active staff posture, current platform counts, F0/F0.5/ MAP, magnet, cost, notification and cron-recovery domains documented; prior baseline classified RED drift

Maintenance Rule

This Markdown file is the living source of truth. PDF copies are snapshots only. A future change is incomplete if it materially changes the product posture, engines, routes, tables, integrations, cost controls, operational jobs or certified-output semantics without updating the affected [ACC-###](#) rows and this update record.