

ATANDA

SPC Player

End-User Guide

Cognitive Production Cockpit (Open Access)

Note: This guide is a standalone documentation PDF explaining how to use the SPC Player. It is completely distinct from the JSON output package generated by the SPC Player itself during a run.

WHAT THE SPC PLAYER IS (AND ISN'T)

What it is: The SPC Player is an open-access cognitive production cockpit. It allows you to register capability briefs, execute them sequentially against a fixed set of Developer Kit cards, evaluate the results through strict governance rubrics, and package the output as a portable JSON artefact. It provides a structured, predictable way to build and review capability execution.

What it is not: It is not a live integration environment. The environment sketches it references are documentation, not active connections. It does not produce certified MVP-PDDs (that remains an ARCHITECT-tier F7 capability), and it does not automatically deploy code.

OPEN ACCESS AND ENTITLEMENTS

The SPC Player is available to all authenticated users.

- No Tier Required: You do not need a paid plan.
- No Billing or Credits: Executions do not consume project credits.
- No SPC-Specific Tier Gate: Normal authentication, provider availability, and platform operating limits still apply.

THE SIX DEV KIT CARDS

You build a run by selecting from the six available Developer Kit cards. They execute strictly in the order you select them.

- ATLAS 360 PLAN: Registered planning card. Its detailed execution contract is outside the supplied SPC Player sources.
- SPARTAN: Scopes the Minimum Viable Product (MVP).
- CODE DJ: Builds the scoped MVP.
- BUGMXT: Debugs the result.
- CODON: Aligns intake against an existing MVP instead of rebuilding it.
- CORDON: Sequences build order.

EXECUTION FLOW & PROFILES

The SPC Player follows a rigid four-step flow:

- **Classify:** Your plain-English brief is deterministically classified (e.g., planning, remediation, review, general).
- **Execute:** The selected cards are invoked in exact order.
- **Evaluate:** The retained results are passed to the governance evaluator.
- **Package:** Everything is assembled into a final JSON output.

Profiles:

- **Full:** Retains all stage evidence and detailed reasoning.
- **Rapid:** Optimises for concise stage content.

Both profiles execute the cards sequentially; the difference is purely in the verbosity of the retained evidence.

The Execution Advisory:

All runs carry a mandatory PRE_BUILD execution advisory. This explicitly marks the output as a user-authorised REVERB v3 derivation, indicating its pre-production nature.

INDEPENDENT GOVERNANCE SCORES

The final evaluation assesses your run across three axes:

- Clarity
- Truthfulness
- Detectability

Crucial Policy: These scores are always reported independently (0 to 100). The SPC Player is strictly prohibited from generating a composite or average score.

RUN HISTORY AND STATUS GLOSSARY

Runs are scoped entirely to you (the owner). You cannot see runs belonging to others, and they cannot see yours.

- **DRAFT:** The brief is registered and cards are selected, but execution has not started.
- **RUNNING:** The engine is actively processing the sequential cards.
- **FAILED:** Execution encountered an error. You can manually retry the run from the dashboard.
- **COMPLETED:** The run finished successfully and the JSON package is ready.

DELIVERY AND SAFE WEBHOOKS

Once a run is COMPLETED, you can download the JSON package directly, or deliver it to an external system via webhook.

Safe Webhook Checklist:

- Must be a fully qualified HTTPS URL.
- Must point to a publicly resolvable IP address (no localhost, 127.0.0.1, internal network addresses, or metadata servers).
- Must not contain embedded credentials (e.g., `https://user:pass@example.com`).

Consent is explicit: You must authorise the endpoint for each specific run before delivery. We do not store secrets or permanent webhook subscriptions.

CURRENT LIMITATIONS

- Not Certified: Runs do not generate public verify URLs (unlike the F7 MVP-PDD process).
- No Automatic Send: Webhook delivery requires explicit manual confirmation per run.
- No Private Targets: Internal VPC or localhost webhooks are deliberately blocked.
- No Credentials: The webhook payload cannot be signed, and custom authentication headers are not supported.
- JSON Only: The SPC Player outputs a JSON package. It does not generate a PDF of the run itself.

STEP-BY-STEP QUICK START

- Register a Brief: Navigate to the SPC Player dashboard and click "Register New Brief". Enter your title, describe the capability, and choose your profile (Full or Rapid).
- Select Cards: Click to select the Dev Kit cards you need. The order you select them is the order they will execute.
- Execute: Click "Execute Run". Watch as the engine progresses through the stages.
- Review: Read the retained stage results and independent governance scores.
- Download or Deliver: Download the final JSON package, or authorise a public HTTPS endpoint and click "Deliver Final Package".

OUTPUT ANATOMY

When you download the JSON package, you will see:

- advisory: The PRE_BUILD execution authority and invoked stages.
- content: The array of retained stage results, each containing a verdict, content, and evidence trail.
- governanceEvaluation: The final textual evaluation of the run.
- scores: The standalone integers for clarity, truthfulness, and detectability.
- distributionPlan: Always records the conservative plan-only policy; an explicitly authorised webhook delivery is a separate user action and does not rewrite that policy.

COMMON QUESTIONS & TROUBLESHOOTING

Q: My execution failed. Do I lose my brief?

A: No. A FAILED run safely retains your brief and card selections. Simply open the run and click "Retry Run".

Q: Why can't I authorise my webhook endpoint?

A: Ensure you are using https://, and verify that the hostname resolves to a public IP. Localhost and private network IPs (like 192.168.x.x or 10.x.x.x) are blocked for security.

Q: Where is my overall score?

A: The SPC Player governance policy explicitly forbids composite scores. You must evaluate the Clarity, Truthfulness, and Detectability metrics independently.