

ATANDA Command Centre

An End-User Guide

Part 1 — Business Cases Catalogue (reality-grounded)

Part 2 — Onboarding via Context Craft: How to design the per-chapter checkpoint pages

Version 1.0 · May 2026

About this guide

The ATANDA Command Centre is, at its core, a small set of high-leverage primitives:

- **An 8-engine HARNESS pipeline** (F1–F8) that takes a raw idea and walks it through diagnostic → atomic prompt → CELL micro-agent → Micro PDD → 15-section SPC → 4-Part ATLAS PDD → SPARTAN-certified MVP PDD → scaffolded code.
- **Adjunct engines** for ingestion, cartridge context injection, vibe-coding tool recommendation (F6-VDJ), drift detection (PFP), and SPC evolution (DE-SPC).
- **A scoring rubric** — JCSE (Junglenomics Composite Score Estimate, 0–50) across 7 pillars, with cert tiers SILVER / GOLD / PLATINUM, calibrated to the HIVE 14-D framework — baked directly into the engine prompts.
- **An exemplar library** users can fork from (SPHINX, SPARTAN, ATLAS, TITAN, GAMEMXT and similar Ultra SI archetypes).
- **A team / organisation layer** with per-seat billing, an org-level activity log, and a cost-cap guard rail.
- **A public verification surface** (/verify?artifactId=..., optionally syndicated to a Sphinx Marketplace) so certified artifacts have an independently checkable URL.

Those primitives compose into more business cases than the platform was originally pitched for. Part 1 of this guide enumerates them. Part 2 then takes a specific, important reuse case — onboarding through the book *Context Craft* — and proposes a concrete chapter-by-chapter checkpoint page design that turns the book into a guided platform onramp.

Everything below is grounded against the actual shipped platform (engines, routes, schema, Stripe SKUs) as of May 2026. Where a use case requires a feature the platform does not have, the guide says so explicitly.

PART 1 — BUSINESS CASES CATALOGUE

The catalogue is organised into ten families. For each business case the guide lists:

- **Who it's for** — the buyer / user persona.
- **What they do on the platform** — the concrete engine sequence.
- **Tier needed** — Explorer / Practitioner / Architect / Team Lite / Institution.
- **Add-ons** — Cartridge, Ingestion, Sphinx publish, etc.

- **What the platform does NOT do** — the boundary; what the buyer still has to do off-platform.
-

FAMILY A — Core builder use cases (what the platform was designed for)

A1. Solo founder building an AI-native product

- **Who:** Technical founder / indie maker building one AI agent product end-to-end.
- **Platform sequence:** F1 (sharpen idea) → F2 (atomic prompt) → F3 (CELL agent) → F4 (Micro PDD) → F5 (15-section SPC, GOLD or PLATINUM) → F6 (ATLAS PDD, 4-Part) → F6-VDJ (pick coding tool) → F7 (SPARTAN MVP PDD with verify URL) → F8 (scaffold code) → PFP (post-F8 drift check).
- **Tier:** Architect (\$199/mo). Architect tier gives unlimited daily F1–F5 and 2 F8 runs per day.
- **Add-ons:** Optional Sphinx publish for a marketing-grade public verify URL.
- **Platform does NOT:** Deploy the agent, train any model, host the production stack, run user analytics. F8 emits a scaffold; turning it into running production is the founder's job.

A2. AI-native startup with a small engineering team

- **Who:** 2–6 person engineering team, every engineer building or extending agent capabilities.
- **Platform sequence:** Architect-tier founder + Team Lite seats for each engineer (per-seat). Team treats each agent capability as its own session: F1 → F5 → F8.
- **Tier:** Architect (founder) + Team Lite seats (engineers).
- **Add-ons:** A Cartridge (\$499.99 one-time) per major product surface to lock the cross-engineer context.
- **Platform does NOT:** Provide a Jira/Linear-style backlog board, sprint planning, or stand-up tooling. The team continues to use whatever issue tracker they already use; the platform's contribution is the certified spec the issue tracker references.

A3. AI agency / boutique consultancy doing client-funded builds

- **Who:** Agency selling AI agent builds to enterprise clients (typical engagement: \$50K–\$500K).
- **Platform sequence:** One session per client engagement. Every deliverable (discovery doc, architecture, build, hand-off) is a HARNESS artifact with a SPARTAN cert ID and a verify URL the client can independently check.
- **Tier:** Architect per consultant + an Institution (per-seat) workspace if running >5 concurrent engagements.
- **Add-ons:** Cartridge per client (loads the client's domain context — security policies, brand guide, existing systems — into every subsequent engine call); Ingestion per legacy spec the client provides.
- **Sells the buyer on:** The verify URL — the client can audit the agency's process without re-reading the agency's deck.
- **Platform does NOT:** Bill the client, manage SOWs, track time. Use Harvest / Bonsai / standard agency tooling.

A4. Enterprise AI transformation programme (brownfield)

- **Who:** F500 with a legacy stack, replatforming workflows onto AI agents.

- **Platform sequence:** Per workstream — Ingestion of the existing spec → Advanced Cartridge with the legacy context → F1 to baseline current prompts/features → F5 to certify replacement SPCs → F8 to scaffold replacements → PFP for the drift report on every PR.
- **Tier:** Institution (per-seat team subs across the programme).
- **Add-ons:** Multiple Cartridges (one per EPIC / workstream); Ingestion per legacy doc.
- **Platform does NOT:** Replace the change-management, parallel-track-migration, SLA-continuity, or org-design work. The platform's contribution is the certified spec + scaffold + drift trail — typically a small fraction of total programme cost.

A5. Greenfield enterprise AI build

- **Who:** Enterprise division building a new AI product (no legacy to migrate).
 - **Platform sequence:** Same as A1 but with Institution-grade per-seat subs, an internal team workspace, and the org activity log as the audit trail.
 - **Tier:** Institution.
 - **Add-ons:** Cartridge per product to lock corporate context (security, brand, integration surfaces).
-

FAMILY B — Education, training, and academic use cases

B1. AI literacy bootcamp (private)

- **Who:** Coding bootcamp / executive education program teaching "how to prompt seriously."
- **Platform sequence:** Cohort of students each gets an Explorer or Practitioner seat. Curriculum maps directly to F1–F7: every assignment is "produce an SPC for problem X with JCSE ≥ 42 ." Instructors review by clicking the artifact's verify URL.
- **Tier:** Explorer (free, F1 + exemplar fork) for sampling; Practitioner for full coursework.
- **Add-ons:** A pre-loaded Cartridge per cohort, containing the cohort's reference materials.
- **Why it works:** JCSE is an objective rubric — students can self-grade, and the score is the same across cohorts. The exemplar library doubles as a teaching corpus.
- **Platform does NOT:** Track attendance, manage cohorts, or issue diplomas. Use standard LMS (Canvas / Moodle) for those.

B2. University AI course / module

- **Who:** Faculty teaching a one-semester AI-product-design course.
- **Platform sequence:** One session per student; pillar-by-pillar weekly assignments scored by JCSE; final project = end-to-end F1→F7 deliverable.
- **Tier:** Practitioner or Explorer for students; one Architect seat for the instructor.
- **Add-ons:** Course Cartridge with reading list and case studies.
- **See also Part 2 of this guide** for a concrete pattern: using the book *Context Craft* as the textbook with per-chapter platform checkpoints.

B3. Corporate L&D — AI fluency programme

- **Who:** Chief Learning Officer rolling out company-wide AI training.

- **Platform sequence:** Each business unit gets an Institution workspace; each learner runs the JST self-assessment, then walks F1 on a real workflow they own. Successful certification = SPC produced for at least one of their daily workflows.
- **Tier:** Institution.
- **Add-ons:** Corporate Cartridge per BU containing internal policy and brand voice.
- **Why it works:** Activity log + per-user /me/costs gives L&D defensible adoption metrics — not "we trained 1,200 people," but "1,200 people produced N certified artifacts with mean JCSE X."
- **Platform does NOT:** Push learning content, schedule sessions, send reminders. Pair with a Cornerstone / Workday Learning / Docebo for delivery.

B4. K-12 / high-school AI literacy

- **Who:** District CTO running a year-long pilot.
- **Platform sequence:** Teacher-led F1 sessions; students score and iterate prompts; final assessment = produce one Micro PDD (F4).
- **Tier:** Explorer for students; Practitioner for teachers.
- **Caution:** Platform has no parental-consent flow, no COPPA hardening, no per-student data-minimisation switches. Districts must wrap their own consent layer.

B5. Train-the-trainer / certification authority

- **Who:** Industry body wanting to define what "AI-fluent" means for their profession (e.g. AI literacy for legal professionals).
- **Platform sequence:** Body designs a Cartridge embodying their profession's domain; certifies trainers by issuing them a verify URL for a body-of-knowledge SPC; trainers then teach using B1/B2 patterns.
- **Tier:** Institution (for the body) + Practitioner (per certified trainer).

B6. AI onboarding for new hires (internal academy)

- **Who:** Engineering org with regular hiring; uses platform as the "your first AI agent in 2 weeks" onboarding track.
- **Platform sequence:** New hire forks an internal exemplar SPC, modifies for a real internal use case, produces a certified MVP PDD, scaffolds with F8, presents to mentor.
- **Tier:** One Team Lite seat per new hire.

FAMILY C — Programmatic & event-based use cases

C1. Hackathon facilitation

- **Who:** Hackathon organiser running a 24–72h event with 50–500 participants.
- **Platform sequence (event-day):**
 - Pre-event: organiser builds one **Cartridge per challenge track** loaded with the prompt brief, judging rubric, sponsor data, and any APIs. Cost: \$499.99 per challenge.
 - Hackers each get an Explorer / Practitioner seat. The Cartridge guarantees every team starts with the *same* authoritative context.

- Judging: every submission is a certified SPC + MVP PDD with a verify URL. Judges score by clicking the URL — no slide decks needed.
- **Tier:** Practitioner per hacker (pay-per-event, sponsor-funded usually); Institution for the organiser.
- **Why it works:** The cartridge-fence guarantee (=== CARTRIDGE CONTEXT (authoritative · do not contradict) === is literally prepended to every engine call) ensures all teams operate from the same brief, eliminating "did you read the rules?" disputes. The verify URL eliminates judging-time confusion.
- **Platform does NOT:** Manage registration, team formation, the leaderboard, or prize disbursement. Use Devpost / HackerEarth for those.

C2. Innovation challenge (internal corporate)

- **Who:** Chief Innovation Officer running a quarterly internal challenge.
- **Platform sequence:** Same as C1 but Cartridge contains internal context. Output: shortlist of certified MVP PDDs the org can actually staff into PIs.
- **Tier:** Institution.

C3. Accelerator / incubator demo-day prep

- **Who:** YC-style accelerator with 10–20 startup cohort.
- **Platform sequence:** Each startup gets one Cartridge (their own market thesis + tech context). They use F1→F7 to produce a public verify URL their demo-day audience can independently audit. Pitch deck links the verify URL.
- **Tier:** Architect per startup; Institution for the accelerator.
- **Add-on:** Sphinx publishing turns each startup's certified MVP PDD into a public-marketplace artifact investors can browse before demo day.

C4. Pitch competition judging

- **Who:** Competition organiser (VC firm running themed pitch contest, university-startup competition, etc.).
- **Platform sequence:** Use JCSE 7-pillar rubric (or a competition-specific Cartridge) as the standardised judging rubric. Each pitch produces an F1 score; finalist shortlist is the top-N JCSE.
- **Tier:** Institution for organiser; Explorer for entrants.
- **Why it works:** JCSE is a published, reproducible rubric — eliminates "judges disagree because they used different criteria" complaints.

C5. Conference workshop / masterclass

- **Who:** Conference organiser hosting a 90-minute hands-on masterclass.
- **Platform sequence:** Workshop Cartridge with the masterclass's worked example loaded; attendees granted 24h Practitioner seats. Output: each attendee leaves with their own verify URL.

FAMILY D — Methodology-as-a-service (the SPC as a universal spec format)

D1. Product management — using ATLAS PDDs as PRDs

- **Who:** PM team standardising on a single rigorous PRD format.
- **Platform sequence:** Every new product surface gets a session; F6 produces the 4-Part ATLAS PDD that replaces the team's previous PRD template.
- **Tier:** Architect per PM.
- **Why it works:** ATLAS PDDs are HIVE-scored — discoverable gaps in the spec (missing constraints, vague success criteria) are flagged automatically.
- **Platform does NOT:** Roadmap, prioritise, or sequence across PRDs. Use Productboard / Aha! for that.

D2. Requirements engineering for non-AI software

- **Who:** Engineering org that wants the rigor of HARNESS specs even for non-AI features.
- **Platform sequence:** Skip F8 (or run it and discard the AI-scaffolded code); use F5 + F6 + F7 as a hardened spec authoring tool.
- **Tier:** Practitioner per spec author.
- **Caveat:** The engines are prompt-tuned for AI agent specs; non-AI specs require user discipline to ignore agent-specific sections.

D3. RFP authoring (buyer side)

- **Who:** Procurement team writing an RFP for a vendor-built system.
- **Platform sequence:** Cartridge with the internal context (existing systems, security policy, regulatory environment); F5 produces a hardened SPC that becomes the RFP. The verify URL becomes a vendor-facing reference artifact.
- **Tier:** Architect or Institution.

D4. RFP response (vendor side)

- **Who:** Vendor responding to AI-system RFPs.
- **Platform sequence:** Ingestion of the RFP; F5 produces a certified counter-spec; the verify URL goes in the proposal as proof-of-rigour.
- **Tier:** Architect per proposal lead.
- **Why it works:** Most RFP responses are PowerPoint slop. A verify URL is differentiating.

D5. Solution architecture documentation

- **Who:** Solution architect at an SI / professional-services firm.
- **Platform sequence:** F6 ATLAS PDD becomes the canonical client-facing architecture document; SPCs (F5) become the per-component specs.
- **Tier:** Architect per SA.

D6. Standards body / certification authoring

- **Who:** Industry body (or an aspiring one) authoring a standard or certification framework.
- **Platform sequence:** The standard itself is an SPC. The certification scheme becomes a Cartridge that downstream practitioners load into their own sessions. The verify URL becomes the publicly auditable canonical reference.
- **Tier:** Institution.

FAMILY E — Regulated and professional services

E1. Legal — drafting standardised agreement templates

- **Who:** Law firm building a template library (e.g. SaaS MSAs, NDAs, employment).
- **Platform sequence:** One Cartridge per practice area (loaded with firm-style examples, jurisdiction notes, partner preferences); F5 produces hardened, JCSE-scored templates; PFP catches drift when the templates are revised.
- **Tier:** Architect per practising attorney; Institution for the firm.
- **Platform does NOT:** Replace privileged review, jurisdiction-specific compliance checking, or the actual practice of law. The output is a draft template, not legal advice.

E2. Healthcare — clinical workflow standardisation

- **Who:** Health system standardising clinician-facing workflow specs (triage protocols, documentation templates).
- **Platform sequence:** Cartridge with the system's clinical guidelines + EHR context; F5 → F6 produces the canonical protocol spec; verify URL = the auditable reference.
- **Tier:** Institution.
- **Platform does NOT:** Carry PHI safely. **Do not** put patient-identifiable data into any engine call. The platform is for protocol specs, not for clinical decisions. No HIPAA hardening ships today.

E3. Compliance / governance policy authoring

- **Who:** CISO / compliance officer writing internal policies (acceptable use, data handling, vendor risk).
- **Platform sequence:** Cartridge per regulatory frame (GDPR / SOC2 / NIST / etc.); F5 produces the policy SPC; F6 the human-readable policy doc.
- **Tier:** Architect or Institution.
- **Platform does NOT:** Maintain a living compliance posture — that's a GRC tool.

E4. Accounting / audit methodology

- **Who:** Audit firm codifying its own playbooks.
- **Platform sequence:** Each audit playbook becomes an SPC; the firm's exemplar SPC library replaces the previous internal wiki.
- **Tier:** Institution.

FAMILY F — Marketing, brand, and editorial

F1. Brand voice system

- **Who:** Marketing org defining and enforcing a brand voice across AI-assisted content production.
- **Platform sequence:** F5 SPC defines the voice (tone, vocabulary, prohibited phrases, audience adaptations); becomes a Cartridge loaded into every downstream content task.
- **Tier:** Architect per brand owner; Institution for cross-team enforcement.

F2. Editorial standards (publishing / newsroom)

- **Who:** Editor-in-chief at a digital publication.
- **Platform sequence:** Style guide → SPC; section guidelines → Cartridges; authors run F1 against their drafts to catch style drift before submission.
- **Tier:** Practitioner per author; Architect per section editor.

F3. Sales playbook standardisation

- **Who:** Sales enablement lead.
- **Platform sequence:** Each "play" (discovery call, demo, closing) gets an SPC; reps use the verify URL as a coachable reference; new-hire ramp uses the Cartridge.
- **Tier:** Institution.

F4. Customer-support agent design

- **Who:** Support ops lead designing agent-handled tier-1 support.
 - **Platform sequence:** F1–F7 to certify the agent design before any vendor build; PFP to monitor drift once deployed.
 - **Tier:** Architect.
-

FAMILY G — Research, science, and academia

G1. Reproducible research methodology cards

- **Who:** Researcher producing methodology supplements for papers.
- **Platform sequence:** Methodology = SPC; verify URL is the reproducibility anchor referenced in the paper.
- **Tier:** Practitioner.

G2. Systematic literature reviews

- **Who:** Meta-research team.
- **Platform sequence:** Cartridge with the corpus; F1 to design the search/screening protocol; F5 to certify the synthesis spec.
- **Tier:** Architect.

G3. Grant proposal authoring

- **Who:** PI writing federal / foundation grants.
- **Platform sequence:** Funder requirements as Ingestion input; F5 SPC of the proposed work; F6 the proposal narrative.
- **Tier:** Architect.

G4. Lab protocol standardisation

- **Who:** Wet lab / engineering lab standardising experimental protocols.
- **Platform sequence:** Each protocol = SPC; the verify URL is referenced in lab notebooks and supplements.
- **Tier:** Practitioner per lab member; Institution for the lab.

FAMILY H — Public sector and civic

H1. Government RFP / RFQ authoring

- **Who:** Government procurement officer.
- **Platform sequence:** Per D3; Cartridge loaded with the relevant procurement code.
- **Tier:** Architect or Institution.
- **Caution:** No FedRAMP / IL4 / GovCloud posture. Public-sector users must keep classified or controlled-unclassified data off-platform.

H2. Policy and regulatory spec drafting

- **Who:** Agency staff drafting rules or guidance.
- **Platform sequence:** Stakeholder input as Ingestion; F5 produces the rule SPC; F6 the human-readable rule.
- **Tier:** Institution.

H3. Civic AI governance frameworks

- **Who:** Municipality designing its own internal AI use policy.
 - **Platform sequence:** Reference frameworks (NIST AI RMF, EU AI Act, OECD) as Ingestion; F5 produces the municipality's own policy SPC.
 - **Tier:** Institution.
-

FAMILY I — Community, marketplace, and IP

I1. Open-source spec libraries

- **Who:** Open-source maintainer publishing reusable SPCs for the community.
- **Platform sequence:** Standard build flow; publish each spec to Sphinx for the public verify URL; community forks into their own sessions.
- **Tier:** Practitioner.

I2. Industry consortium standards

- **Who:** Consortium maintaining a shared standard library.
- **Platform sequence:** Same as I1 but with consortium governance overlay (which member orgs can amend which SPCs is a governance question the platform does not enforce — manage via your consortium's own bylaws).
- **Tier:** Institution.

I3. Indie creators monetising prompt IP

- **Who:** Prompt designer / consultant who wants to sell hardened prompts.

- **Platform sequence:** F5 → F6 → F7 produces a certified SPC. Sphinx publish makes it discoverable. Buyers fork into their own session.
- **Tier:** Architect.
- **Why it works:** A SPARTAN cert ID and verify URL is the closest thing to a "this prompt has been peer-reviewed" badge that currently exists.

14. Sphinx Marketplace publishing (revenue-share path)

- **Who:** Any creator with publishable certified artifacts.
- **Platform sequence:** Build → publish to Sphinx → marketplace handles discovery, licensing, payment (where configured).
- **Caveat:** The publish endpoint exists in the platform; the marketplace itself is external and only live where SPHINX_BASE_URL points somewhere real.

FAMILY J — Internal operations and HR

J1. Interview rubric and question bank

- **Who:** Hiring manager / talent ops.
- **Platform sequence:** F5 SPC per role profile (rubric + question bank + scorecard); HR uses the same SPC across interviewers; PFP detects rubric drift over time.
- **Tier:** Architect per hiring manager; Institution for HR ops.

J2. Vendor / tool evaluation templates

- **Who:** Procurement or engineering ops doing structured vendor selection.
- **Platform sequence:** SPC defines the eval criteria; each vendor proposal Ingested and scored against it.
- **Tier:** Architect.

J3. Runbook / SOP standardisation

- **Who:** Ops / SRE team.
- **Platform sequence:** Each runbook = SPC. Verify URL becomes the canonical reference linked from incident channels.
- **Tier:** Architect per team lead; Team Lite per operator.

J4. Knowledge-base authoring

- **Who:** Tech writer / support ops.
- **Platform sequence:** Each KB article = certified SPC, ensuring tone and completeness pass JCSE before publication.
- **Tier:** Practitioner.

Cross-cutting capabilities that unlock these use cases

Capability	Why it matters across families
Verify URL	Makes any output independently auditable — works for clients, judges, students, regulators.
Cartridge context	Lets a team enforce a single source of authoritative context across many users / sessions. Universally useful.
JCSE 7-pillar rubric	Gives every use case a standardised quality score, so quality conversations stop being subjective.
Activity log + <code>/me/costs</code>	Defensible adoption + spend metrics for any organisational rollout.
Cost-cap guard rail (402 <code>COST_CAP_EXCEEDED</code>)	Lets finance say yes to broad rollouts without runaway-spend risk.
Exemplar library	Lowers the cold-start cost for every use case above.

What the platform does NOT do (boundary line, applies across all families)

- Manage projects, sprints, backlogs, time, or money flows.
- Host, deploy, observe, or operate the resulting AI agents.
- Push notifications at cost-cap thresholds (it 402s at the cap, no warnings before).
- Carry PHI / classified / export-controlled data with regulatory hardening.
- Replace change-management, procurement, legal review, clinical judgment, or scientific peer review.
- Issue diplomas, certificates of attendance, or accredited credentials (verify URLs are not regulatory certifications).
- Run as a multi-agent runtime panel — engine calls are one at a time, one user at a time; "Ultra SI" personas are archetypes in prompts and forks in the exemplar library, not parallel runtime agents.

PART 2 — ONBOARDING VIA CONTEXT CRAFT: HOW TO DESIGN THE PER-CHAPTER CHECKPOINT PAGES

The book *Context Craft: The Last Human Skill — The Rise of Cognitive Engineering* is a near-perfect platform onboarding vehicle. Its 7-pillar pedagogy (SYSTEM, ROLE, INSTRUCTION, DATA, CONSTRAINTS, STRUCTURE, VOICE/STANDARD) maps directly onto the 7 pillars that F1's JCSE diagnostic scores, and the book's Chapter 8 "Super Prompt" integration corresponds exactly to running F2 → F5 on the platform for the first time. Chapter 9's economic argument (the \$21M discrepancy) corresponds to the `/me/costs` dashboard. Chapter 10+ self-assessment corresponds to the JST self-rating.

The reader of *Context Craft* who finishes the book has already learned everything they need to be productive on the platform — they just don't yet know that the platform is the place where the doctrine becomes

operational. A well-designed checkpoint page after each chapter turns the book from "interesting read" into "ten-chapter guided onramp" with the platform as the natural next click.

This section proposes a concrete checkpoint design.

The book's existing structure (what each chapter already gives you)

Each chapter already follows a 4-part scaffold:

- **Part 1 — STORYLINE** (the narrative anchor; David, Marcus, Victoria, etc.)
- **Part 2 — SAMPLE PROMPTS** (Before & After Architecture)
- **Part 3 — CARD ACTIVATION** (the pillar in context)
- **Part 4 — QUEST CARD** (activate the pillar)
- **CHAPTER BRIDGE — SYNERGY LOOP** (transition to the next chapter)

Add a fifth element — **CHECKPOINT** — between QUEST CARD and CHAPTER BRIDGE. Keep it to exactly one PDF page. The reader's eye is already trained to expect a beat at the end of each chapter; the checkpoint slots in cleanly.

Single-page checkpoint template (use the same template across all chapters)

[illegible]

Key design rules behind the template:

- 1 **One page, no overflow.** A checkpoint that spills onto a second page breaks the chapter rhythm.
- 2 **5-minute platform task ceiling.** If the reader can't complete the checkpoint in five minutes, fewer will do it. Larger tasks belong in Chapter 8+.
- 3 **Always the same five sections** (recap · do-this · score · stuck? · what's next). Predictability builds confidence.
- 4 **The score box is the hook.** Readers will fill in scores compulsively if you give them the box. By Chapter 7 they will *want* to know their cumulative JCSE.

- 5 **Every checkpoint produces an artifact with a verify URL** — that URL is the reader's portable proof they did the work. By Chapter 8 they have seven of them.

Chapter-by-chapter content map

Ch	Pillar	Platform action	Tier needed	Exemplar to fork	Target score
1	SYSTEM	Sign up; fork SYSTEM exemplar; run F1; capture SYSTEM pillar score	Explorer (free)	"SPHINX ULTRA SI" (rich SYSTEM definition)	5/7
2	ROLE	Fork ROLE exemplar; run F1; ROLE pillar score	Explorer	"ATLAS ULTRA SI" (sharp role anchor)	5/7
3	INSTRUCTION	Same flow, INSTRUCTION pillar	Explorer	"SPARTAN" (instructional precision)	5/7
4	DATA	Fork an exemplar and upload one document as your first mini-Cartridge (\$0 dry-run — let the reader see the cartridge fence in their own session)	Practitioner	"TITAN ULTRA SI"	5/7
5	CONSTRAINTS	Run F1 on a deliberately under-constrained prompt; observe the pillar flag	Practitioner	"GAMEMXT ULTRA SI"	4/7 (this is the hardest pillar; calibrate the bar lower)
6	STRUCTURE	Run F4 (Micro PDD) on the prompt you've been building	Practitioner	Any chapter exemplar	5/7

Ch	Pillar	Platform action	Tier needed	Exemplar to fork	Target score
7	VOICE / STANDARD	Run F1 on your prompt with a deliberate "house voice" instruction; observe the VOICE pillar score	Practitioner	"Editorial voice" exemplar (purpose-built for the book)	5/7
8	INTEGRATION (Super Prompt)	First full F2 → F3 → F5 run. The reader produces their first 15-section SPC and a SILVER or GOLD cert. This is the moment to convert Explorer → Architect.	Architect (\$199/mo)	Their own work	JCSE total ≥ 35/50
9	ECONOMICS	Tour /me/costs; run F7 to produce a SPARTAN-certified MVP PDD with a verify URL; share the URL with one peer	Architect	—	Cert SILVER+
10	SELF (JST)	Run the JST self-assessment (in-book worksheet; results feed into the user's exemplar library annotation)	—	—	JST baseline captured
11	TEAM	(If the book covers team adoption) Create an org; invite one teammate; run a session together	Team Lite seat	—	First org-visible session

Ch	Pillar	Platform action	Tier needed	Exemplar to fork	Target score
12	CAPSTONE	F8 Code DJ on the SPC built in Chapter 8; PFP for drift; readers leaves with running scaffolded code	Architect	—	F8 produced + PFP green

The "Workbook" companion (recommended pull-out)

Bind a 16-page perforated workbook at the back of the book containing:

- A **JCSE running-score sheet** (one box per chapter, plus a "cumulative" running total).
- A **JST self-assessment** with cut-out scoring grid (tied to Chapter 10).
- A **session-log table** the reader fills in as they go (chapter · session ID · verify URL · date · score).
- A final **certification page** the reader can scan and submit to claim a "Context Craft Practitioner" badge on the platform (this requires a small platform feature — see "Platform changes to support this onboarding" below).

The workbook is the physical artefact that links the book to the platform: when readers complete it, they have demonstrably run the whole pipeline.

Voice and tone of the checkpoints (match the book)

The book's voice is direct, sometimes confrontational, anchored in concrete dollar-impact stories (David's \$21M, Marcus's \$200M). Match it:

- **Direct second-person.** "Go to the platform. Run F1. Look at your score." Not "users can optionally navigate to..."
- **Stakes-anchored.** "The CONSTRAINTS pillar is where Marcus lost \$23 million. You're going to score yourself on it in the next four minutes."
- **Honest about difficulty.** "Most readers score 3/7 on STRUCTURE the first time. So did David. Run it again."
- **No platform jargon in the checkpoint body.** "F1" is fine (the book teaches the F-numbers); but no "endpoint" / "schema" / "Stripe" — those leak the implementation and break the reader's flow.

Pacing and reader friction

Empirical pattern from similar book-to-platform onramps:

- **Chapters 1–3:** ~40% of readers will do the checkpoints in real time. Keep these the easiest possible (Explorer tier only; pure fork-and-run).
- **Chapters 4–7:** Drops to ~20% completing in real time. This is fine — the workbook captures asynchronous completion.
- **Chapter 8:** This is the conversion event (Explorer → Architect \$199/mo). The checkpoint must feel like a meaningful milestone — getting the first SILVER cert. Design accordingly: bigger box, optional callout, "your

first certified artifact." Consider a Practitioner trial offer (\$49/mo, 14-day) that ramps to Architect after Chapter 8.

- **Chapter 12:** This is the moment to ask for a referral / community signup. The reader has produced running code. They will share if asked.

Platform changes recommended to support this onboarding

These are small platform additions that would make the book-to-platform flow materially stronger:

- 1 **"Onboarding tracks" feature.** A track ID (`context_craft_v1`) the user can claim on signup that pre-loads the right exemplar at each F1 run and auto-tallies the per-pillar score in their dashboard. Without this, readers track manually in the workbook — works, but loses the dashboard pleasure.
- 2 **A "claim my book progress" page.** Reader uploads a photo of the completed workbook → platform marks the user as a "Context Craft Practitioner" and surfaces a badge on their public profile / verify URLs. Cheap to build, high emotional payoff.
- 3 **A book-specific Cartridge.** A free Cartridge containing the seven worked examples from the book, available to any reader. Eliminates the cold-start for Chapter 1.
- 4 **A book-cohort discount SKU.** Promo code printed inside the book back cover → 50% off Architect for the first three months. Conversion lever for Chapter 8.
- 5 **A per-pillar exemplar library** explicitly curated for teaching. The current exemplar library is curated for builders; for learners, you want each exemplar to be a high-leverage demonstration of one pillar at near-perfect score, with the other pillars deliberately neutral.

These five changes are each ~1–2 weeks of platform work. None require new engines.

What success looks like

If the checkpoint pages are designed as above:

- **Read-through completion** (reader finishes Chapter 12): held by the book's existing pull, not changed.
- **Platform signup conversion** (reader creates an account): targets ~20–30% of readers — the Chapter 1 checkpoint is the gateway.
- **Pillar-score reporting** (reader completes ≥ one F1 run): targets ~15–20%.
- **Paid conversion at Chapter 8** (Explorer/Practitioner → Architect): targets ~5–8% of total readers, ~25–35% of those who reached Chapter 8 on-platform.
- **F8-run completion** (reader produces scaffolded code via Chapter 12 capstone): targets ~2–3% of total readers — small but extremely high-value, because these readers become the platform's reference customers.

The numbers are achievable because the book and the platform are doing complementary work: the book is teaching the *why* and the *what*; the checkpoint pages convert each chapter's *what* into a five-minute hands-on *how*; the platform delivers the proof-of-work artifact (verify URL) that closes the loop. Without the checkpoints, the book is a great book; with them, it becomes the platform's most effective onboarding asset and a durable revenue lever for years.

End of guide.

Generated against ATANDA Command Centre platform state, May 2026, and against the Context Craft: The Last Human Skill manuscript (12 chapters, prologue + 2 epilogues).